



**Figure 9: Performance comparison of FGD and baselines
(Monte-Carlo Workload Inflation on Alibaba Trace)**


(a) Unallocatable GPUs given arriving workloads



(b) GPU nodes occupied during scheduling



(c) GPU requests of failed tasks (at 96% capacity)



(d) Fragmentation breakdown into three causes

